

Database Tuning - Introduction

Vu Tuyet Trinh

trinhvt@soict.hust.edu.vn

Department of Information Systems
SoICT-HUST

Copyrights:

Many slides belong to the tutorial:

Database Tuning

Principles, Experiments and Troubleshooting Techniques

Dennis Shasha (shasha@cs.nyu.edu)

Philippe Bonnet (bonnet@diku.dk)

and some from the web ...



What is Database Tuning?

Activity of making a database application run faster:

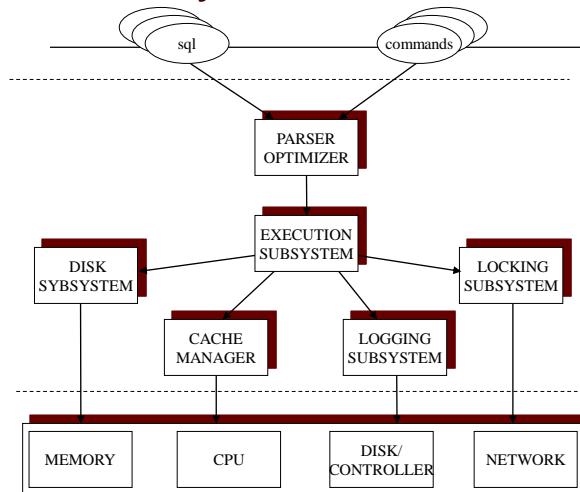
- Faster means higher throughput (or response time)
- Avoiding transactions that create bottlenecks or avoiding queries that run for hours unnecessarily is a must.
- A 5% improvement is significant.



Why Database Tuning?

- Troubleshooting:
 - Make managers and users happy given an application and a DBMS
- Capacity Sizing:
 - Buy the right DBMS given application requirements
- Application Programming:
 - Coding your application for performance

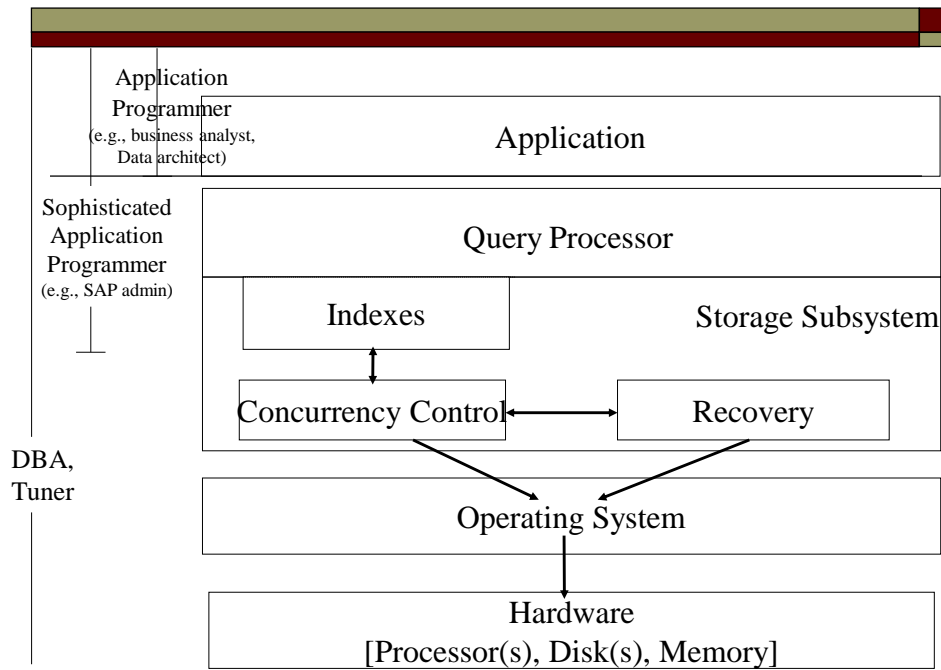
Why is Database Tuning hard?



The following query runs too slowly

```
select *
from R
where R.a > 5;
```

What do you do?





Outline

- ❑ Schema tuning
- ❑ Index tuning
- ❑ Query Processing
- ❑ Query tuning
- ❑ Transaction Management
- ❑ Transaction tuning
- ❑ Tuning Distributed Application



Tuning Principles

1. Think globally, fix locally
2. Partitioning breaks bottlenecks
 - temporal and spatial
3. Start-up costs are high; running costs are low
4. Render unto server what is due unto server
5. Be prepared for trade-offs



Think globally, fix locally

- Proper identification of problem; minimal intervention
- Understand the whole, including the application goals before taking a set of queries and find the indexes that speed them up.
- **Example:**
 - High I/O, paging and processor utilization may be due to frequent query scans instead of using an index or log sharing a disk with some frequently accessed data.



Partitioning breaks bottlenecks

- Technique for reducing the load on a certain component of the system either by dividing the load over more resources or by spreading the load over time
- Partitioning may not always solve bottleneck:
 - First, try to speed up the component
 - If it doesn't work, partition
- **Example:**
 - Lock and resource contention among long and short transactions

Start-up costs are high; running costs are low

- Obtain the effect you want with the fewest possible start-ups
- Examples:
 - It is expensive to begin a read operation on a disk, but once it starts disk can deliver data at high speed.
 - So, frequently scanned tables should be laid out consecutively on disk.
 - Cost of parsing, semantic analysis, and selecting access paths for simple queries is significant
 - So, often executed queries should be compiled

Render unto server what is due unto server

- Important design question is the allocation of work between the DB system (server) and the application program (client)
- Depends on:
 - Relative computing resources of client and server
 - Where the relevant information is located
 - Whether the DB task interacts with the screen

Be prepared for trade-offs

- Increasing speed of application requires combination of memory, disk and computational resources
- Examples:
 - Adding an index => speeds up critical query, but increases disk storage, and space in RAM
 - Increasing RAM => Decreasing I/O, speed up query, but spending more money

