

Index tuning

Vu Tuyet Trinh

trinhvt@soict.hust.edu.vn

Department of Information Systems
SoICT-HUST

Copyrights:

Many slides belong to the tutorial:

Database Tuning

Principles, Experiments and Troubleshooting Techniques

Dennis Shasha (shasha@cs.nyu.edu)

Philippe Bonnet (bonnet@diku.dk)


and some from the web ...

File Organization

- Data storage in file
 - records, blocks and access structures
- Organisation
 - **Heap files:** for full file scans or frequent updates
 - Data unordered
 - Write new data at end
 - **Index:** if retrieved in key attributes
 - Need to store index
 - **Sorted Files:** if retrieved in sort order or want range
 - Need external sort or an index to keep sorted
 - **Hashed Files:** if selection on equality
 - Collection of buckets with primary & overflow pages
 - Hashing function over search key attributes

Heap File


- Organization
 - Data unordered
 - Write new data at end



Need a full scan file for Search, Insert, Update, Delete operations

Index

An index is a data structure that supports efficient access to data




Classes of Indexes

- **Primary vs. secondary:** primary has primary key
- **Clustered vs. unclustered:** order of records and index approximately same
 - Alternative 1 implies clustered, but not vice-versa
 - A file can be clustered on at most one search key
- **Dense vs. Sparse:** dense has index entry per data value; sparse may “skip” some
 - Alternative 1 always leads to dense index
 - Every sparse index is clustered!
 - Sparse indexes are smaller; however, some useful optimizations are based on dense indexes

B+-Tree

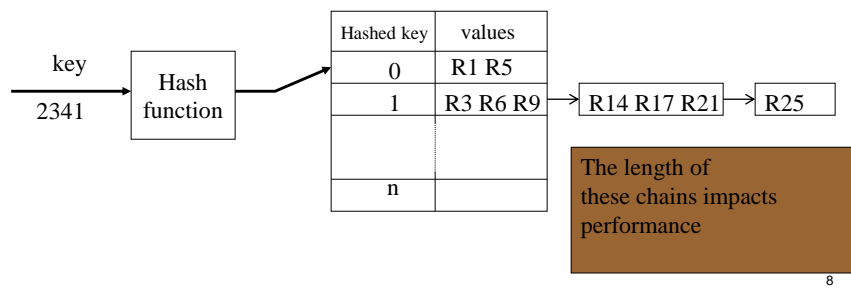
- A B+-Tree is a balanced tree whose leaves contain a sequence of key-pointer pairs.



7

Hash Index

- A hash index stores key-value pairs based on a pseudo-randomizing function called a *hash function*.




8

Clustered vs. Unclustered Index

Suppose Index Alternative (2) used, records are stored in Heap file


- Perhaps initially sort data file, leave some gaps
- Inserts may require overflow pages



Dense / Sparse Index


□ Sparse index

- Pointers are associated to pages



□ Dense index

- Pointers are associated to records
- Non clustered indexes are dense



Index Implementations in some major DBMS

- SQL Server
 - B+-Tree data structure
 - Clustered indexes are sparse
 - Indexes maintained as updates/insertions/deletions are performed
- DB2
 - B+-Tree data structure, spatial extender for R-tree
 - Clustered indexes are dense
 - Explicit command for index reorganization
- Oracle
 - B+-tree, hash, bitmap, spatial extender for R-Tree
 - No clustered index until 10g
 - Index organized table (unique/clustered)
 - Clusters used when creating tables.
- MySQL
 - B+-Tree, R-Tree (geometry and pairs of integers)
 - Indexes maintained as updates/insertions/deletions are performed¹¹

Constraints and Indexes

- Primary Key, Unique
 - A non-clustered index is constructed on the attribute(s) that compose the primary key with the constraint that values are unique.
- Foreign Key
 - By default, no index is created to enforce a foreign key constraint.



Performance Issues

- ❑ Type of Query
- ❑ Index Overhead

13



Types of Queries

1. Point Query

```
SELECT balance
FROM accounts
WHERE number = 1023;
```

3. Range Query

```
SELECT number
FROM accounts
WHERE balance > 10000;
```

2. Multipoint Query

```
SELECT balance
FROM accounts
WHERE branchnum = 100;
```

4. Prefix Match Query

```
SELECT *
FROM employees
WHERE name = 'Jensen'
      and firstname = 'Carl'
      and age < 30;
```

14

Types of Queries

5. Extremal Query

```
SELECT *  
FROM accounts  
WHERE balance =  
    max(select balance from  
    accounts)
```

6. Ordering Query

```
SELECT *  
FROM accounts  
ORDER BY balance;
```

7. Grouping Query

```
SELECT branchnum,  
    avg(balance)  
FROM accounts  
GROUP BY branchnum;
```


8. Join Query

```
SELECT distinct branch.adresse  
FROM accounts, branch  
WHERE  
    accounts.branchnum =  
        branch.number  
and accounts.balance > 10000;
```

15

Index

An index is a data structure that supports efficient access to data



16



Index Tuning Knobs

- ❑ Index data structure
- ❑ Search key
- ❑ Size of key
- ❑ Clustered/Non-clustered/No index
- ❑ Covering

17



Clustered Index

Benefits of a clustered index:

1. A sparse clustered index stores fewer pointers than a dense index.
 - This might save up to one level in the B-tree index.
2. A clustered index is good for multipoint queries
 - White pages in a paper telephone book
3. A clustered index based on a B-Tree supports range, prefix, extremal and ordering queries well.

18



Clustered Index

4. A clustered index (on attribute X) can reduce lock contention:
Retrieval of records or update operations using an equality, a prefix match or a range condition based on X will access and lock only a few consecutive pages of data

Cost of a clustered index

1. Cost of overflow pages
 - Due to insertions
 - Due to updates (e.g., a NULL value by a long string)

19



Clustered Index

- Because there is only one clustered index per table, it might be a good idea to replicate a table in order to use a clustered index on two different attributes
 - Yellow and white pages in a paper telephone book
 - Low insertion/update rate

20

Non-Clustered Index

Benefits of non-clustered indexes

1. A dense index can eliminate the need to access the underlying table through covering.
 - It might be worth creating several indexes to increase the likelihood that the optimizer can find a covering index

2. A non-clustered index is good if each query retrieves significantly fewer records than there are pages in the table.

- Point queries

- Multipoint queries:

*number of distinct key values >
 $c * \text{number of records per page}$*

Where c is the number of pages retrieved in each prefetch

21

Index on Small Tables

- Tuning manuals suggest to avoid indexes on small tables
 - If all data from a relation fits in one page then an index page adds an I/O
 - If each record fits in a page then an index helps performance



Key Compression

- Use key compression
 - If you are using a B-tree
 - Compressing the key will reduce the number of levels in the tree
 - The system is not CPU-bound
 - Updates are relatively rare


23



Summary

1. Use a hash index for point queries
Use a B-tree for multipoint and range queries
2. Use clustering
 - if your queries need all or most of the fields of each records returned
 - if range queries are asked
3. Use a dense index to cover critical queries
4. Don't use an index if the time lost when inserting and updating overwhelms the time saved when querying

24



25