Query Processing

Vu Tuyet Trinh

trinhvt@soict.hust.edu.vn

Department of Information Systems SoICT-HUST

Parser Query plan Optimized Code Generator

Code for executing

Outline

- □ Standard relational operators
- □ Query plans & exec strategies
- Query Optimization

Relational Algebra

- Set of operations on relations
- □ Basic of the implementation of most current database systems.
- Operation
 - Input: one or two relation
 - Output: a relation
- □ Relational algebra expression or algebraic expression
 - a sequence of relational algebra operations

Relational Algebra Operations

- □ Projection
- □ Selection
- □ Join
- Division
- Union
- □ Intersection
- Difference
- Cartesian product

-

Example

Student

010010111		
ld	Name	Suburb
1108	Robert	Kew
3936	Glen	Bundoor a
8507	Norman	Bundoor a
8452	Mary	Balwyn

Student_Sport

ld	Name	Suburb	SportID	Sport
1108	Robert	Kew	05	Swimmin g
3936	Glen	Bundoor a	05	Swimmin g
8507	Norman	Bundoor a	05	Swimmin g
8452	Mary	Balwyn	05	Swimmin g
1108	Robert	Kew	09	Dancing
3936	Glen	Bundoor a	09	Dancing
8507	Norman	Bundoor a	09	Dancing
8452	Mary	Balwyn	09	Dancing

Sport



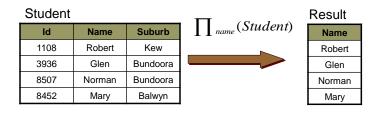
SportID	Sport
05	Swimming
09	Dancing

Projection

□ Choose some attributes.

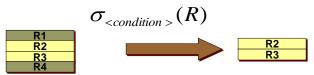


Example



Selection

□ choose from R each tuple where the condition holds.



* Example

Student

ld	Name	Suburb
1108	Robert	Kew
3936	Glen	Bundoora
8507	Norman	Bundoora
8452	Mary	Balwyn

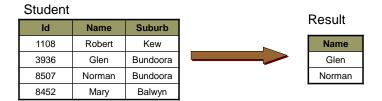


 $\sigma_{\textit{suburb}="Bundoora}(\textit{Student})$

F	Result		
I	ld	Name	Suburb
I	3936	Glen	Bundoora
I	8507	Norman	Bundoora

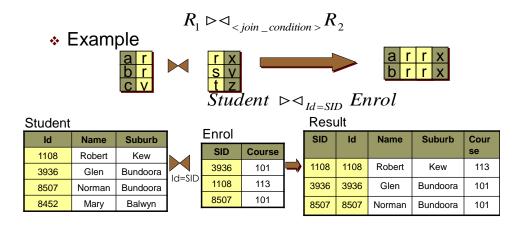
Example

$$\prod_{name} (\sigma_{suburb="Bundoora} Student)$$



Join

□ Combine attributes from 2 tables



5

Example

$$\prod_{\textit{name}, \textit{Course}} (\sigma_{\textit{suburb} = "\textit{Bundoo}}(\textit{Student} \ \rhd \lhd_{\textit{Id} = \textit{SID}} \textit{Enrol}\))$$

Student

ld	Name	Suburb
1108	Robert	Kew
3936	Glen	Bundoora
8507	Norman	Bundoora
8452	Mary	Balwyn

Enrol

SID	Course
3936	101
1108	113
8507	101



Cours e
101
101

11

Natural Join

□ Special join operation with equal join condition on their common attributes

*

□ Example

lakes	
SID	SNO
1108	21
1108	23
8507	23
8507	29

Enrol

SID	Course
3936	101
1108	113
8507	101



SID	SNO	Course
1108	21	113
1108	23	113
8507	23	101
8507	29	101

Outer-Join

□ Left Outer Join



□ Right Outer Join



Example

Student

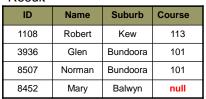
ID	Name	Suburb
1108	Robert	Kew
3936	Glen	Bundoora
8507	Norman	Bundoora
8452	Mary	Balwyn



Enrol

SID	Course
3936	101
1108	113
8507	101

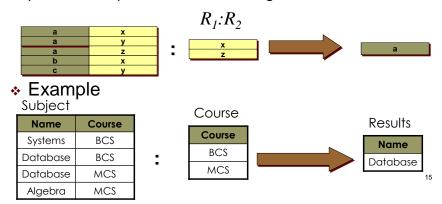
Result





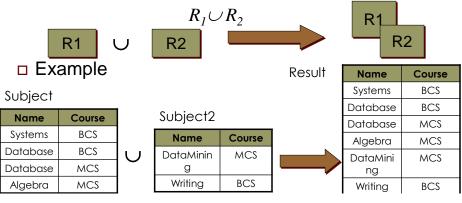
Division

□ The division operator divides a dividend relation R1 or degree m+n by a divisor relation R2 of degree n, and produces a quotient relation of degree m.



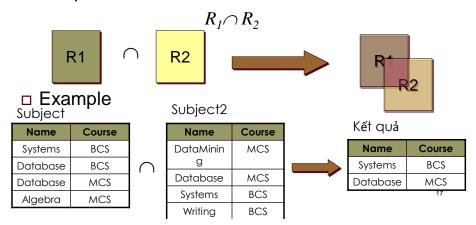
Union

 combining the tuples from two input unioncompatible relations (having the same set of attributes)



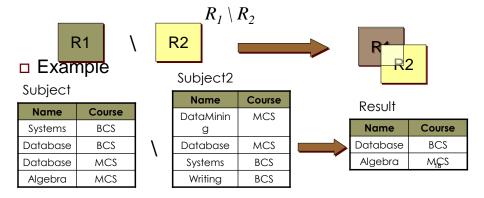
Intersection

□ Keeping only common tuples from 2 input unioncompatible relation



Difference

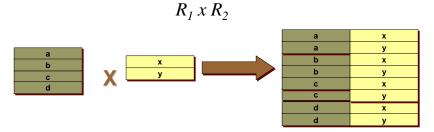
□ Difference of two union compatible relations is a relation containing tuples occurred in the first relation but not in the second.



9

Cartesian Product

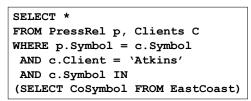
□ the concatenation of every tuple of one relation with every tuple of the other relation.

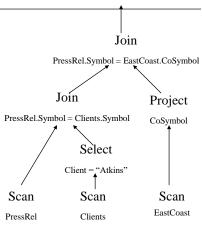


19

Query Plans

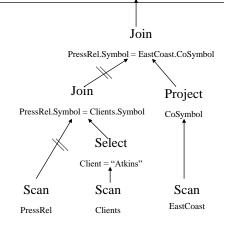
- Data-flow graph of relational algebra operators
- □ *Typically:* determined by optimizer





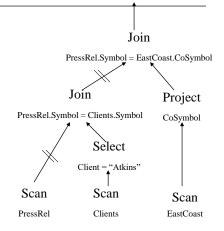
Execution Strategy Issues

- ☐ Granularity & parallelism:
 - Pipelining vs. blocking
 - Materialization



Iterator-Based Query Execution

- Execution begins at root
 - open, next, close
 - Propagate calls to children May call multiple child nexts
- Efficient scheduling & resource usage



Basic Principles

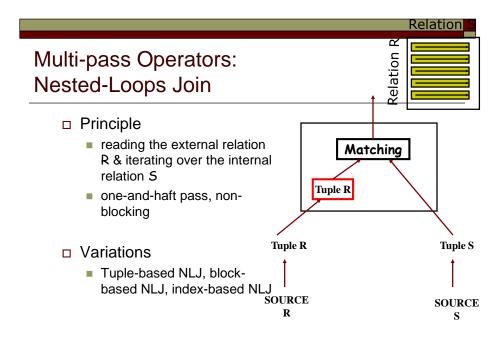
- Many DB operations require reading tuples, tuple vs. previous tuples, or tuples vs. tuples in another table
- □ Techniques generally used:
 - Iteration: for/while loop comparing with all tuples on disk
 - Index: if comparison of attribute that's indexed, look up matches in index & return those
 - Sort: iteration against presorted data (interesting orders)
 - Hash: build hash table of the tuple list, probe the hash table
- Must be able to support larger-than-memory data

Basic Operators

- □ One-pass operators:
 - Scan
 - Select
 - Project
- Multi-pass operators:
 - Join
 - Various implementations
 - □ Handling of larger-than-memory sources
 - Semi-join
 - Aggregation, union, etc.

1-Pass Operators: Scanning a Table

- □ Sequential scan: read through blocks of table
- □ Index scan: retrieve tuples in index order
- □ Cost in page reads -- b(T) blocks, r(T) tuples
 - b(T) pages for sequential scan
 - Up to r(T) for index scan if unclustered index
 - Requires memory for one block



Two-Pass Algorithms

Sort-based

Need to do a multiway sort first (or have an index) Approximately linear in practice, 2 b(T) for table T

Hash-based

Store one relation in a hash table

(Sort-)Merge Join

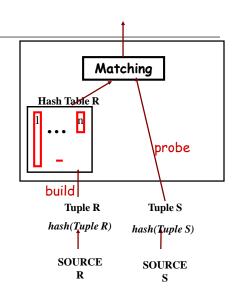
- □ Requires data sorted by join attributes
 - Merge and join sorted files, reading sequentially a block at a time
 - Maintain two file pointers
 - □ While tuple at R < tuple at S, advance R (and vice versa)
 - □ While tuples match, output all possible pairings
 - Preserves sorted order of "outer" relation
- Very efficient for presorted data
- Can be "hybridized" with NL Join for range joins
- May require a sort before (adds cost + delay)
- □ Cost: b(R) + b(S) plus sort costs, if necessary In practice, approximately linear, 3 (b(R) + b(S))

Hash-Based Joins

- □ Allows partial pipelining of operations with equality comparisons
- □ Sort-based operations block, but allow range and inequality comparisons
- Hash joins usually done with static number of hash buckets
 - Generally have fairly long chains at each bucket
 - What happens when memory is too small?

Hash Join (HJ)

- Principle
 - using a hash table (HT) of R
 - two-pass, blocking algorithm
- Process
 - build the HT of R
 - probe every tuple of S with corresponding tuples in the HT of R

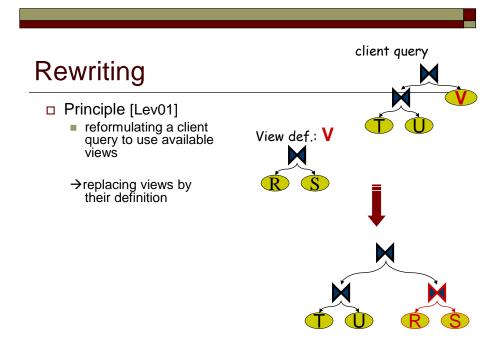


Overview of Query Optimization

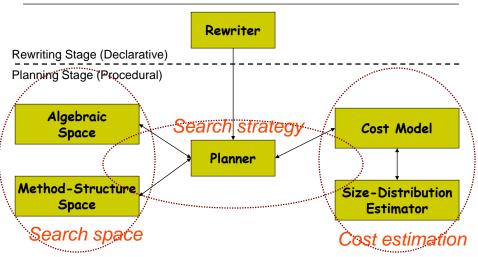
- □ Plan: Tree of R.A. ops, with choice of alg for each op.
 - Each operator typically implemented using a `pull' interface: when an operator is `pulled' for the next output tuples, it `pulls' on its inputs and computes them.
- □ Two main issues:
 - For a given query, what plans are considered?
 - □ Algorithm to search plan space for cheapest (estimated) plan.
 - How is the cost of a plan estimated?
- □ Ideally: Want to find best plan.

Practically: Avoid worst plans!

Query optimization [loa00] Rewriter Planning Stage (Declarative) Planning Stage (Procedural) Algebraic Space Planner Size-Distribution Estimator



Query optimization [loa00]



Search Space

- □ Algebraic space
 - operators execution orders that are to be considered by the Planner for each query sent to it
- □ Method-structure space
 - implementation choices that exist for the execution of each specified ordered series of operators

Relational Algebra Equivalences

□ Allowing to choose different join orders and to `push' selections and projections ahead of joins.

Selections:
$$\sigma_{c1 \wedge ... \wedge cn}(R) \equiv \sigma_{c1}(\ldots \sigma_{cn}(R))$$
 (Cascade) $\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c2}(\sigma_{c1}(R))$ (Commute)

• Projections:
$$\pi_{a1}(R) \equiv \pi_{a1}(...(\pi_{an}(R)))$$
 (Cascade)

* Joins:
$$R \bowtie (S \bowtie T) \equiv (R \bowtie S) \bowtie T$$
 (Associative)
 $(R \bowtie S) \equiv (S \bowtie R)$ (Commute)

Show that: $R \bowtie (S \bowtie T) \equiv (T \bowtie R) \bowtie S$

More Equivalences

- □ A projection commutes with a selection that only uses attributes retained by the projection.
- □ Selection between attributes of the two arguments of a cross-product converts cross-product to a join.
- □ A selection on just attributes of R commutes with join (i.e., σ (R ⋈ S) ≡ σ (R) ⋈ S)
- □ Similarly, if a projection follows a join R ⋈ S, we can `push' it by retaining only attributes of R (and S) that are needed for the join or are kept by the projection.

Enumeration of Alternative Plans

- □ There are two main cases:
 - Single-relation plans
 - Multiple-relation plans
- □ For queries over a single relation, queries consist of a combination of selects, projects, and aggregate ops:
 - Each available access path (file scan / index) is considered, and the one with the least estimated cost is chosen.
 - The different operations are essentially carried out together (e.g., if an index is used for a selection, projection is done for each retrieved tuple, and the resulting tuples are *pipelined* into the aggregate computation).

Cost estimation

- Cost factors
 - CPU cost + I/O cost + Communication cost
 - Size of (intermediate) relations
- □ For each plan considered, must estimate cost:
 - Must estimate cost of each operation in plan tree.
 - Depends on input cardinalities.
 - Must also estimate size of result for each operation in tree!
 - Use information about the input relations.
 - □ For selections and joins, assume independence of predicates.

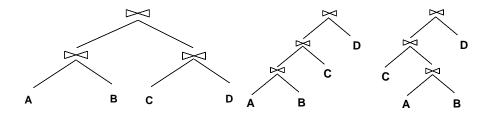
Search strategy

- □ Objective
 - exploring the set of alternative execution plans and finding the "cheapest" one
- □ Taxonomy
 - Polynomial vs. Combinatorial
 - Heuristics vs. Systematic
 - Deterministic vs. Randomized
 - Transformative vs. Constructive

orthogonal but related

Queries Over Multiple Relations

- Fundamental decision in System R: <u>only left-deep join trees</u> are considered.
 - As the number of joins increases, the number of alternative plans grows rapidly; we need to restrict the search space.
 - Left-deep trees allow us to generate all fully pipelined plans.
 - □ Intermediate results not written to temporary files.
 - □ Not all left-deep trees are fully pipelined (e.g., SM join).

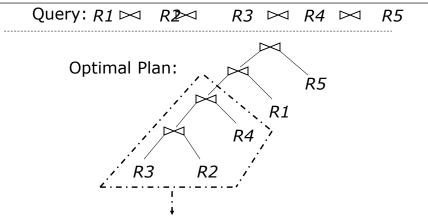


Enumeration of Left-Deep Plans

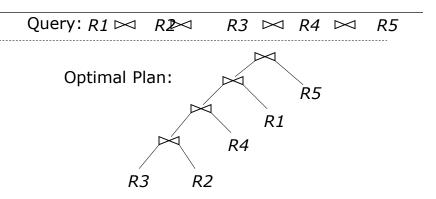
- □ Left-deep plans differ only in the order of relations, the access method for each relation, and the join method
- Enumerated using N passes (if N relations joined):
 - Pass 1: Find best 1-relation plan for each relation.
 - Pass 2: Find best way to join result of each 1-relation plan (as outer) to another relation. (All 2-relation plans.)
 - Pass N: Find best way to join result of a (N-1)-relation plan (as outer) to the N'th relation. (All N-relation plans.)
- □ For each subset of relations, retain only:
 - Cheapest plan overall, plus
 - Cheapest plan for each interesting order of the tuples.

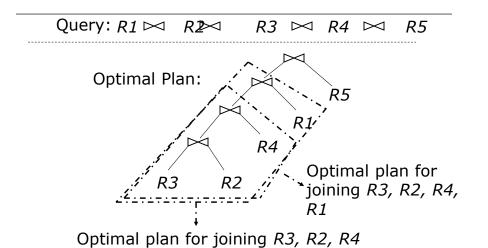
Enumeration of Plans (Contd.)

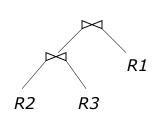
- ORDER BY, GROUP BY, aggregates etc. handled as a final step, using either an 'interestingly ordered' plan or an addional sorting operator.
- □ An N-1 way plan is not combined with an additional relation unless there is a join condition between them, unless all predicates in WHERE have been used up.
 - i.e., avoid Cartesian products if possible.
- ☐ In spite of pruning plan space, this approach is still exponential in the # of tables.



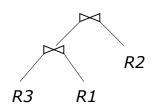
Optimal plan for joining R3, R2, R4



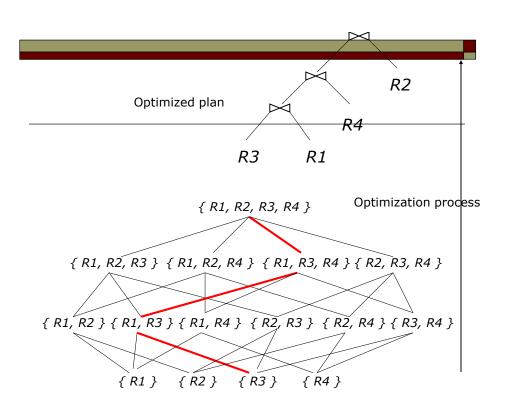




Optimal for joining *R1, R2, R3*



Sub-Optimal for joining *R1*, *R2*, *R3*



Performance issues

- Disk access
- Index utilization
- Database schema
- Nested queries, nested transaction
- □ ...

Example

- □ Employee(<u>ssnum</u>, name, manager, dept, salary, numfriends)
 - Clustering index : ssnum
 - Non clustering indexes (i) name (ii) dept (iii) salary
- □ Student(**ssnum**, name, degree_sought, year)
 - Clustering index :ssnum
 - Non clustering index :name
- □ Tech(<u>dept</u>, manager, location)
 - Clustering index : dept

Rewriting queries

- Using Index
- Eliminating DISTINCTs
- Nested queries
- Join queries
- Having clause
- Using view
- Using materialized views

Using index

- Unenabling the use of index
 - Math expression WHERE salary/12 >= 4000;
 - String function SELECT * FROM employee WHERE SUBSTR(name, 1, 1) = 'G';
 - Comparison with different data type
 - Comparison with null value

DISTINCTs

■ Example

SELECT distinct ssnum FROM employee WHERE dept = 'information systems'

■ Necessary?

Example

SELECT DISTINCT ssnum FROM employee, tech WHERE employee.dept = tech.dept

■ DISTINCT or not?

Employee(<u>ssnum</u>, name, manager, dept, salary, numfriends)
Student(<u>ssnum</u>, name, degree_sought, year)
Tech(<u>dept</u>, manager, location)

55

- SELECT ssnum
 FROM employee, tech
 WHERE employee.manager = tech.manager
- SELECT ssnum, tech.dept
 FROM employee, tech
 WHERE employee.manager = tech.manager
- SELECT student.ssnum
 FROM student, employee, tech
 WHERE student.ssnum = employee.ssnum
 AND employee.dept = tech.dept;

Nested queries

- □ SELECT ssnum FROM employee WHERE salary > (select avg(salary) from employee)
- SELECT ssnum FROM employee
 WHERE dept in (select dept from tech)
- SELECT ssnum
 FROM employee e1
 WHERE salary =
 (SELECT avg(e2.salary)
 FROM employee e2,tech
 WHERE e2.dept = e1.dept AND
 e2.dept = tech.dept)

Nested queries

SELECT ssnum
FROM employee
WHERE dept in (select dept from tech)

→

SELECT ssnum FROM employee, tech WHERE employee.dept = tech.dept

Nested queries with aggregation function

- Example
 - SELECT avg(salary)
 FROM employee
 WHERE manager in (select manager from tech)
 - SELECT avg(salary)
 FROM employee, tech
 WHERE employee.manager = tech.manager
- □ Any difference ?

Nested queries vs. join queries

Example

SELECT ssnum
FROM employee e1
WHERE salary = (SELECT avg(e2.salary
FROM employee e2, tech
WHERE e2.dept = e1.dept
and e2.dept = tech.dept);

- Create view temp as SELECT avg(salary) as avsalary, employee.dept FROM employee, tech WHERE employee.dept = tech.dept GROUP BY employee.dept;
- SELECT ssnum
 FROM employee, temp
 WHERE salary = avsalary
 AND employee.dept = temp.dept

Employee(<u>ssnum</u>, name, manager, dept, salary, numfriends) Student(<u>ssnum</u>, name, degree_sought, year) Tech(**dept**, manager, location)

SELECT ssnum

FROM employee e1

WHERE numfriends = COUNT(SELECT e2.ssnum

FROM employee e2, tech

WHERE e2.dept = tech.dept

AND e2.dept = e1.dept);

- INSERT INTO temp SELECT COUNT(ssnum) as numcolleagues, employee.dept FROM employee, tech WHERE employee.dept = tech.dept GROUP BY employee.dept;
- SELECT ssnum
 FROM employee, temp
 WHERE numfriends = numcolleagues
 AND employee.dept = temp.dept;

Join condition

- □ Join over clustering indexes.
- □ Join expression vs. data type (numeric, string,...)

Using view

- □ CREATE VIEW techlocation
 AS
 SELECT ssnum, tech.dept,
 location
 FROM employee, tech
 WHERE
 employee.dept =
 tech.dept;
- SELECT location
 FROM techlocation
 WHERE ssnum = 43253265;
- SELECT location
 FROM employee, tech
 WHERE
 employee.dept =
 tech.dept
 AND ssnum =
 43253265;

Materialized view

Materialized views:

CREATE MATERIALIZED VIEW VendorOutstanding BUILD IMMEDIATE REFRESH COMPLETE ENABLE QUERY REWRITE AS SELECT orders.vendor, sum(orders.quantity*item.price)

FROM orders, item
WHERE orders.itemnum = item.itemnum
group by orders.vendor;

Parameters

- BUILD immediate/deferred
- REFRESH complete/fast
- ENABLE QUERY REWRITE

Advantages

- Automatically update
- Used by optimizer

