


Transaction Management

Vu Tuyet Trinh

trinhvt@soict.hust.edu.vn


Department of Information Systems
SoICT-HUST

Transaction – example



Transaction

- A sequence of read and write operations on data items that logically functions as one unit of work
 - Assuring data integrity and correction
- ACID Properties
 - Atomicity
 - Consistency
 - Isolation
 - Durability



3

Atomicity

- guarantee that either all of the tasks of a transaction are performed or none of them are

- Example

```
T: Read(A,t1);  
  If t1 > 500 {  
    Read(B,t2);  
    t2:=t2+500;  
    Write(B,t2);  
    t1:=t1-500;  
    Write(A,t1);  
  }
```

A diagram showing a transaction flow. Four horizontal arrows point to the right, corresponding to the four statements in the transaction code: 'Read(A,t1);', 'Write(B,t2);', 't1:=t1-500;', and 'Write(A,t1);'. The second arrow, corresponding to 'Write(B,t2);', is crossed out with a large 'X' and the word 'crash' is written next to it.

4

Consistency

- ensures that the DB remains in a consistent state before the start of the transaction and after the transaction is over
- Example

```
T: Read(A,t1);  
  If t1 > 500 {  
    Read(B,t2);  
    t2:=t2+500;  
    Write(B,t2);  
    t1:=t1-500;  
    Write(A,t1);  
  }
```

← $A+B = C$

← $A+B = C$

5

Isolation

- ability of the application to make operations in a transaction appear isolated from all other operations.
- Example $A= 5000, B= 3000$

```
T: Read(A,t1);  
  If t1 > 500 {  
    Read(B,t2);  
    t2:=t2+500;  
    Write(B,t2);  
    t1:=t1-500;  
    Write(A,t1);  
  }
```

← $T': A+B$
 $(= 5000+3500)$

← $(A+B = 4500+3500)$

6


Durability

- guarantee that once the user has been notified of success, the transaction will persist, and not be undone
- Ví dụ: $A = 5000$, $B = 3000$

```
T: Read(A,t1);  
  If t1 > 500 {  
    Read(B,t2);  
    t2:=t2+500;  
    Write(B,t2);  
    t1:=t1-500;  
    Write(A,t1);  
  }  
                                ← crash  
                                A= 4500, B=3500
```

7

Transaction States



8

Transaction Management Interfaces

- Begin Trans
- Commit ()
- Abort()

- Savepoint Save()
- Rollback (savepoint)
(savepoint = 0 ==> Abort)

9

Concurrency Control

□ Objective:

- ensures that database transactions are performed concurrently without the concurrency violating the data integrity
- guarantees that no effect of committed transactions is lost, and no effect of aborted (rolled back) transactions remains in the related database.

□ Example

<pre>T0: read(A); A := A - 50; write(A); read(B); B := B + 50; write(B);</pre>	<pre>T1: read(A); temp := A * 0.1; A := A - temp; write(A); read(B); B := B + temp; write(B);</pre>
--	---

10

Scheduling

T ₀	T ₁	T ₀	T ₁	T ₀	T ₁
read(A) A := A - 50 write(A) read(B) B := B + 50 write(B)	read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B)	read(A) A := A - 50 write(A) read(B) B := B + 50 write(B)	read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B)	read(A) A := A - 50 write(A) read(B) B := B + 50 write(B)	read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B)
(1)		(2)		(3)	11

Serializability

- A **schedule** of a set of transactions is a linear ordering of their actions
 - e.g. for the simultaneous deposits example:

R1(X) R2(X) W1(X) W2(X)
- A **serial schedule** is one in which all the steps of each transaction occur consecutively
- A **serializable schedule** is one which is equivalent to some serial schedule

Lock

□ Definition

- a synchronization mechanism for enforcing limits on access to DB in concurrent way.
- one way of enforcing concurrency control policies

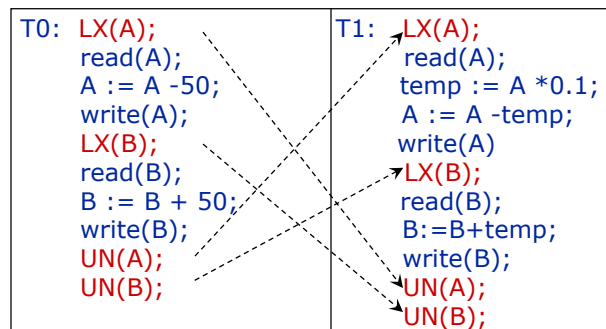
□ Lock types

- **Shared lock (LS)** readable but can not write
- **Exclusive lock (LX)**: read and write
- UN(D): unlock

□ Compatibility

	LS	LX
LS	true	false
LX	false	false

Example




Well-Formed, two-phased transaction

- A transaction is **well-formed** if it acquires at least a shared lock on Q before reading Q or an exclusive lock on Q before writing Q and doesn't release the lock until the action is performed
 - Locks are also released by the end of the transaction
- A transaction is **two-phased** if it never acquires a lock after unlocking one
 - i.e., there are two phases: a *growing phase* in which the transaction acquires locks, and a *shrinking phase* in which locks are released

2Phase Locking (2PL)

- Phase 1
 - locks are acquired and no locks are released
- Phase 2
 - locks are released and no locks are acquired



Example

T ₁		T ₃	
Lock(A)		Lock(B)	
Read(A)		Read(B)	
Lock(B)		B=B-50	
Read(B)		Write(B)	
B:=B+A		Unlock(B)	
Write(B)		Lock(A)	
Unlock(A)		Read(A)	
Unlock(B)		A=A+50	
		Write(A)	
		Unlock(A)	
	T ₂		T ₄
	Lock(B)		Lock(A)
	Read(B)		Read(A)
	Lock(A)		Unlock(A)
	Read(A)		Lock(B)
	Unlock(B)		Read(B)
	A:=A+B		Unlock(B)
	Write(A)		Printn(A+B)
	Unlock(A)		

2PL

Not 2PL

Deadlock

T0: LX(B); (1)	T1: LX(A); (4)
read(B); (2)	read(A); (5)
B := B + 50; (3)	temp := A * 0.1; (6)
write(B); (8)	A := A - temp; (7)
LX(A); (10)	write(A) (9)
read(A);	LX(B);
A := A - 50;	read(B);
write(A);	B:=B+temp;
UN(A);	write(B);
UN(B);	UN(A);
	UN(B);



Resolving Deadlock

- Detecting
 - Recovery when deadlock happen
 - rollback
 - Used waiting-graph
- Avoiding
 - Resource ordering
 - Timeout
 - Wait-die
 - Wound-wait



Waiting Graph

- Graph
 - Node handling lock or waiting for lock
 - Edge $T \rightarrow U$
 - U handle L(A)
 - T wait to lock A
 - T must wait until U unlock A
- If there exists a cycle in the waiting graph \rightarrow deadlok



Timeout


- ❑ Set a limit time for each transaction
- ❑ If time-out → do rollback



Exercises

Transaction – consistency

collection of action that preserve consistency



with assumption


IF T starts with consistent state +
 T executes in isolation
THEN T leaves consistent state

How can constraints be violated?

- ❑ Transaction bug
- ❑ DBMS bug
- ❑ Hardware failure
 e.g., disk crash
- ❑ Data sharing
 e.g., T1 and T2 in parallel

Failures

Events — Desired
 \ Undesired — Expected
 \ Unexpected




Recovery

- ❑ Maintaining the consistency of DB by ROLLBACK to the last consistency state.
- ❑ Ensuring 2 properties
 - Atomic
 - Durability
- Using LOG




Transaction Log

- A sequence of log record keeping trace of actions executed by DBMS
 - <start T>
 - Log the beginning of the transaction execution
 - <commit T>
 - transaction is already finished
 - <abort T>
 - Transaction is calcel
 - <T, X, v, w>
 - Transaction makes an update actio, before update $X=v$, after update $x = w$

- 
-
- Read(A)
 - If $A > 50$ then display(“so du hop le”)
 - Else {
 - $A:=A+50$
 - =====→CRASH
 - display (“ghi no tai khoan A”)
 - }

Transaction Log

- Handled in main memory and put to external memory (disk) when possible



Checkpoint

- Definition:
 - moment where intermediate results and a log record are saved to disk.
 - being initiated at specified intervals
- Objective
 - minimize the amount of time and effort wasted when restart
 - the process can be restarted from the latest checkpoint rather than from the beginning.
- Log record
 <checkpoint> or <ckpt>

Undo-logging

Step	Action	t	Mem A	Mem B	Disk A	Disk B	Mem Log
1							<start T>
2	Read(A,t)	8	8		8	8	
3	t:=t*2	16	8		8	8	
4	Write(A,t)	16	16		8	8	<T, A, 8>
5	Read(B,t)	8	16	8	8	8	
6	t:=t*2	16	16	8	8	8	
7	Write(B,t)	16	16	16	8	8	<T, B, 8>
8	Flush log						
9	Output(A)	16	16	16	16	8	
10	Output(B)	16	16	16	16	16	
11							<commit T>
12	Flush log						

Undo-Logging Rules

- (1) For every action generate undo log record (containing old value)
- (2) Before X is modified on disk, log records pertaining to X must be on disk (write ahead logging: WAL)
- (3) Before commit is flushed to log, all writes of transaction must be reflected on disk

Undo Logging Recovery Rules

- Let S is set of unfinished transactions
 - $\langle \text{start } T_i \rangle$ in log
 - $\langle \text{commit } T_i \rangle$ or $\langle \text{abort } T_i \rangle$ is not in log
- For each $\langle T_i, X, v \rangle$ in log
 - If $T_i \in S$ then
 - Write(X, v)
 - Output(X)
- For each $T_i \in S$
 - Write $\langle \text{abort } T_i \rangle$ to log

Undo-Logging & Checkpoint

```
<start T1>
<T1, A, 5>
<start T2>
<T2, B, 10>
<T2, C, 15>
<T2, D, 20>
<commit T1>
<commit T2>
<checkpoint>
<start T3>
<T3, E, 25>
<T3, F, 30>
```

scan

```
<start T1>
<T1, A, 5>
<start T2>
<T2, B, 10>
<start ckpt (T1, T2)>
<T2, C, 15>
<start T3>
<T1, D, 20>
<commit T1>
<T3, E, 25>
<commit T2>
<end ckpt>
<T3, F, 30>
```

scan

Redo-logging

Step	Action	t	Mem A	Mem B	Disk A	Disk B	Mem Log
1							<start T>
2	Read(A,t)	8	8		8	8	
3	t:=t*2	16	8		8	8	
4	Write(A,t)	16	16		8	8	<T, A, 16>
5	Read(B,t)	8	16	8	8	8	
6	t:=t*2	16	16	8	8	8	
7	Write(B,t)	16	16	16	8	8	<T, B, 16>
8							<commit T>
9	Flush log						
10	Output(A)	16	16	16	16	8	
11	Output(B)	16	16	16	16	16	
							<T, end>


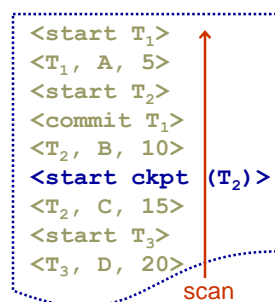
Redo-logging Rules

- (1) For every action, generate redo log record (containing new value)
- (2) Before X is modified on disk (DB), all log records for transaction that modified X (including commit) must be on disk
- (3) Flush log at commit
- (4) Write END record after DB updates flushed to disk

Redo-logging Recovery Rules

- Let S = set of transactions with
 - $\langle T_i, \text{commit} \rangle$ in log
 - no $\langle T_i, \text{end} \rangle$ in log
- For each $\langle T_i, X, w \rangle$ in log, in forward order (earliest \rightarrow latest)
 - If $T_i \in S$ then write(X, w)
output(X)
- For each $T_i \in S$
 - write $\langle T_i, \text{end} \rangle$

Redo Logging & Checkpoint



Discussion

□ Undo Logging

- need to write to disk as soon transaction finishes
- Access disk

□ Redo Logging

- need to keep all modified blocks in memory until commit
- Use memory

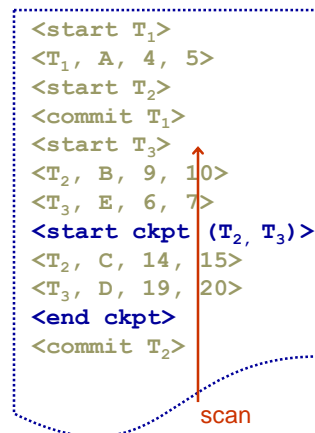

Undo/Redo Loggin

Step	Action	t	Mem A	Mem B	Disk A	Disk B	Mem Log
1							<start T>
2	Read(A,t)	8	8		8	8	
3	t:=t*2	16	8		8	8	
4	Write(A,t)	16	16		8	8	<T, A, 8, 16>
5	Read(B,t)	8	16	8	8	8	
6	t:=t*2	16	16	8	8	8	
7	Write(B,t)	16	16	16	8	8	<T, B, 8, 16>
8	Flush log						
9	Output(A)	16	16	16	16	8	
10							<commit T>
11	Output(B)	16	16	16	16	16	

Undo/Redo Logging Rules

- ❑ Page X can be flushed before or after T commit
- ❑ Log record flushed before corresponding updated page (WAL)
- ❑ Flush at commit (log only)

Undo/Redo Logging & Checkpoint



Undo/Redo Logging Recovery Rules

- ❑ Backwards pass (end of log → latest valid checkpoint start)
- ❑ Constructing set S of committed transactions
- ❑ undo actions of transactions not in S
- ❑ undo pending transactions
- ❑ follow undo chains for transactions in (checkpoint active list) – S
- ❑ Forward pass (latest checkpoint start → end of log)
- ❑ redo actions of S transactions





Isolation Levels

- ❑ Read Uncommitted (No lost update)
 - Exclusive locks for write operations are held for the duration of the transactions
 - No locks for read
- ❑ Read Committed (No inconsistent retrieval)
 - Shared locks are released as soon as the read operation terminates.
- ❑ Repeatable Read (no unrepeatable reads)
 - Strict two phase locking
- ❑ Serializable (no phantoms)
 - Table locking or index locking to avoid phantoms